

Contract Governance Copilot

Governed policy retrieval for contract and procurement teams, delivered in Microsoft Teams

Author	Alex Kinyua
Platform	Microsoft Azure, Microsoft Foundry, Microsoft Copilot Studio
Delivery channel	Microsoft Teams and Microsoft 365 Copilot
Status	Built, validated and published
Source & documentation	github.com/A1-lex/governed-contract-governance-copilot

Summary

Procurement and contract teams lose time to a recurring question: what does our policy actually say? Answering it usually means interrupting Legal, searching a document library, or relying on memory. The first is expensive, the second is slow, and the third is risky.

The Contract Governance Copilot answers those questions inside Microsoft Teams using only the organisation's approved governance documents. Every answer names the document it came from. When a question falls outside the approved knowledge, the assistant says so and notifies the document owner rather than guessing.

The result is that policy becomes self-service, decisions become traceable to a named document, and knowledge gaps surface as actionable notifications instead of silent assumptions.

1. What the assistant does

The assistant is scoped to a single, well-defined domain: contract and procurement governance. It draws on seven approved policy documents covering approvals, procurement, legal escalation, supplier onboarding, contract review, renewal and termination, and conflicts of interest.

Typical questions it handles

- What approvals are required for an NDA?
- When does Legal have to review a contract?
- Who owns supplier onboarding?
- What is the renewal notice window for contracts above \$250,000?
- What is the gift limit for meals and hospitality?

Four behaviours define the product

Behaviour	What it means in practice
Grounded	Answers come only from the approved document set. The assistant will not answer from general knowledge, even when it could.
Cited	Every answer ends by naming the source document, so a reviewer can verify the claim against the policy itself.
Honest about gaps	When the answer is not in the approved documents, the assistant says so rather than producing a plausible guess.
Actionable	An unanswered question becomes an email to the document owner naming the missing topic, turning a gap into a task.

2. Worked example

A question the assistant can answer

Reviewer asks:

“What approvals are required for NDAs?”

The assistant returns the NDA approval thresholds, distinguishing standard-template NDAs, which a department manager may approve, from non-standard or amended NDAs, which require head of department approval and mandatory Legal review. It closes with:

Source: Approval Matrix.pdf; Legal Escalation Guidelines.pdf

A question the assistant cannot answer

Reviewer asks:

“What is our parental leave policy?”

The assistant replies:

"That information isn't in my connected knowledge. Would you like me to notify the content owner?"

On confirmation, the document owner receives an email naming the missing topic and asking whether the policy should be added or the reviewer directed elsewhere. The refusal is not a dead end. It is the start of a content-governance workflow.

3. Why not a general-purpose assistant

A general assistant will answer a policy question confidently whether or not it knows the organisation's actual policy. In a governance context that is worse than no answer, because a plausible response invites action.

This assistant differs in four ways that matter to a compliance function.

	General assistant	Contract Governance Copilot
Source of answers	Model training data and, optionally, the open web	Only the organisation's approved governance documents
Unknown questions	Produces a plausible answer	Declines and offers to notify the document owner
Traceability	No reliable link to an authoritative document	Every answer names its source document
Staying current	Depends on retraining or manual context	New policies are indexed hourly with no change to the assistant

The refusal is the product. An assistant that declines to answer a question outside its approved knowledge is one a Legal team can endorse. An assistant that always answers is one they must police.

4. How it works

Documents are stored in Azure Blob Storage. An Azure AI Search indexer runs hourly, converting new or updated documents into a searchable, vectorised index. A Microsoft Foundry knowledge base sits over that index and performs agentic retrieval: it decomposes a question into sub-queries, searches in parallel, reranks the results, and synthesises an answer with source references. Copilot Studio calls that knowledge base through a workflow and presents the result in Teams.

Component responsibilities

Component	Responsibility
Azure Blob Storage	Holds the approved governance documents. Adding a policy means placing a file here, with no change to the assistant.
Azure AI Search	Chunks, embeds and indexes the documents on an hourly schedule, so newly added policies become answerable without redeployment.
Microsoft Foundry	Hosts the knowledge base that performs query planning, parallel retrieval, semantic reranking and answer synthesis with citations.
Power Automate	Calls the knowledge base retrieval endpoint and returns grounded results with source references to the assistant.
Copilot Studio	Owns the reviewer experience: instructions, grounding rules, citation format, escalation behaviour and Teams delivery.
Office 365 Outlook	Sends the knowledge-gap notification to a fixed, pre-configured document owner.

Request flow

```
graph TD; A[Reviewer question in Microsoft Teams] --> B[Copilot Studio agent]; B --> C[Retrieval workflow]; C --> D[Foundry knowledge base → query planning and reranking]; D --> E[Azure AI Search index ← hourly ingestion from Blob Storage]; E --> F[Grounded answer with source citation]; F --> G[If no answer found → email to document owner];
```

Reviewer question in Microsoft Teams
↓
Copilot Studio agent
↓
Retrieval workflow
↓
Foundry knowledge base → query planning and reranking
↓
Azure AI Search index ← hourly ingestion from Blob Storage
↓
Grounded answer with source citation
↓
If no answer found → email to document owner

5. Evidence

The following screenshots record the assistant's behaviour against the criteria set out in section 1.




Figure 1. A grounded answer to an NDA approval question, closing with a named source document.




Figure 2. The assistant declines an out-of-domain question rather than answering from general knowledge.




Figure 3. A question outside the approved knowledge produces an offer to notify the document owner.




Figure 4. The document owner receives an email naming the missing topic and requesting the source document.




Figure 5. Retrieval validation showing the correct source document returned for each benchmark question.




Figure 6. The assistant published and available to reviewers in Microsoft Teams.

6. Validation

A benchmark set maps each test question to the document that should answer it. Retrieval is judged correct when the expected document is returned as the leading source.

Test question	Expected source
What approvals are required for NDAs?	Approval Matrix
When should Legal review a contract?	Legal Escalation Guidelines
Who owns supplier onboarding?	Supplier Onboarding Process
What documents are required before procurement approval?	Procurement Policy
What is the standard contract review workflow?	Contract Review Procedure
Who approves early termination above \$250,000?	Contract Renewal and Termination Policy
What is the gift limit for meals and hospitality?	Conflict of Interest and Gifts Policy

Behavioural tests

Scenario	Required behaviour
In-scope question	Grounded answer ending with a named source document
Out-of-domain question	Declines to answer, with no response drawn from general knowledge
In-scope topic, absent from documents	States the gap and offers to notify the document owner
Notification confirmed	Email sent to the fixed recipient naming the missing topic
Newly added document	Answerable after the next scheduled ingestion, with no change to the assistant

7. Known limitations

These are stated plainly because they shape what the assistant should and should not be relied on for.

Limitation	Detail and mitigation
Answer quality is bounded by document quality	The assistant is only as accurate as the approved documents. Outdated or conflicting policy produces outdated or conflicting answers. Document ownership and review dates are therefore part of the operating model, not an afterthought.
Grounding is enforced by instruction	The assistant is instructed to answer only from retrieved content. This holds consistently in testing but is a strong constraint rather than a hard architectural guarantee. Behaviour should be re-tested after any model or instruction change.
Escalation recipient must be fixed	Early testing showed the assistant substituting the signed-in user's address for the configured recipient. The recipient is now locked in configuration and reinforced by instruction, because a user-redirectable notification is a governance risk.

Limitation	Detail and mitigation
No per-user document filtering	All reviewers see the same approved document set. Document-level permission trimming is a production-hardening item for deployments where policy visibility varies by role.
Retrieval uses a preview interface	Agentic retrieval is called through a preview API, which may change. The underlying index is generally available and provides a stable fallback path if required.

8. Roadmap

Enhancement	Value
Permission-aware retrieval	Restrict answers to documents the individual reviewer is entitled to see, enabling use across functions with differing access.
Knowledge-gap reporting	Aggregate unanswered questions into a view of where policy coverage is thin, turning escalations into a content roadmap.
Usage and confidence signals	Track which documents answer the most questions and where reviewers disagree with the response.
Broader document estate	Extend beyond contract governance to adjacent policy domains using the same ingestion and citation pattern.

9. Demonstration

A short walkthrough covering the problem, a live demonstration and the architecture behind it.

https://drive.google.com/file/d/1LddBr5te_b10TLfhTF7Y9KBHJvQzpsa7/view?usp=sharing

This assistant is the first of two related builds. The second, the Contract Intelligence Copilot, applies the same governance documents to supplier agreements: where this assistant answers what the policy says, the second assesses whether a given contract complies with it.